

The code is nicely written and it was fun to observe the actions of the agent during its attempts to reach the terminal state.

Being able to watch the learned actions is a nice addition to look more into the agent's behaviors.

Overall, everything works. The code is easy to follow and you covered everything necessary. Good job all!

Additionally I ran one of the notebooks first (Homework 2.ipynb), and it was interesting to choose the actions by myself :D